

```
import pandas as pd
df=pd.read_csv('household_power_consumption.csv')
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity
0	16/12/2006	17:24:00	4.216	0.418	234.84	
1	16/12/2006	17:25:00	5.36	0.436	233.63	
2	16/12/2006	17:26:00	5.374	0.498	233.29	
3	16/12/2006	17:27:00	5.388	0.502	233.74	
4	16/12/2006	17:28:00	3.666	0.528	235.68	

```
df.dtypes
```

→	Date	object
	Time	object
	Global_active_power	object
	Global_reactive_power	object
	Voltage	object
	Global_intensity	object
	Sub_metering_1	object
	Sub_metering_2	object
	Sub_metering_3	float64
	dtype:	object

```
df.shape #(rows,columns)
```

→	(1048575, 9)
---	--------------

```
df['Sub_metering_3'].astype('object')
```

→	0	17.0
	1	16.0
	2	17.0
	3	17.0
	4	17.0
		...
	1048570	0.0
	1048571	0.0
	1048572	0.0
	1048573	0.0
	1048574	0.0
	Name:	Sub_metering_3, Length: 1048575, dtype: object

```
df.info()
```

→	<class 'pandas.core.frame.DataFrame'>
	RangeIndex: 1048575 entries, 0 to 1048574
	Data columns (total 9 columns):
	# Column Non-Null Count Dtype
	--- ---
0	Date 1048575 non-null object
1	Time 1048575 non-null object
2	Global_active_power 1048575 non-null object
3	Global_reactive_power 1048575 non-null object
4	Voltage 1048575 non-null object
5	Global_intensity 1048575 non-null object
6	Sub_metering_1 1048575 non-null object
7	Sub_metering_2 1048575 non-null object
8	Sub_metering_3 1044506 non-null float64
	dtypes: float64(1), object(8)
	memory usage: 72.0+ MB

```
df.dtypes
```

→	Date	object
	Time	object
	Global_active_power	object
	Global_reactive_power	object
	Voltage	object
	Global_intensity	object
	Sub_metering_1	object
	Sub_metering_2	object
	Sub_metering_3	float64
	dtype:	object

```
df.values
```

```
array([['16/12/2006', '17:24:00', '4.216', ..., '0', '1', 17.0],
       ['16/12/2006', '17:25:00', '5.36', ..., '0', '1', 16.0],
       ['16/12/2006', '17:26:00', '5.374', ..., '0', '2', 17.0],
       ...,
       ['13/12/2008', '21:36:00', '0.422', ..., '0', '0', 0.0],
       ['13/12/2008', '21:37:00', '0.422', ..., '0', '0', 0.0],
       ['13/12/2008', '21:38:00', '0.422', ..., '0', '0', 0.0]],
      dtype=object)
```

```
df1=df.sort_values('Global_active_power',ascending=True)
df1.head(1000)
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Globa
894447	28/8/2008	20:51:00	0.076		0	234.03
894448	28/8/2008	20:52:00	0.076		0	233.92
894446	28/8/2008	20:50:00	0.076		0	234.06
894445	28/8/2008	20:49:00	0.076		0	234.34
894444	28/8/2008	20:48:00	0.076		0	234.88
...
883702	21/8/2008	9:46:00	0.08		0	240.71
867366	10/8/2008	1:30:00	0.08		0	241.05
867367	10/8/2008	1:31:00	0.08		0	241.01
884885	22/8/2008	5:29:00	0.08		0	239.63
884886	22/8/2008	5:30:00	0.08		0	239.47

1000 rows × 9 columns

```
df1=df.sort_values('Global_active_power',ascending=False)#descending order
df1.head(4000)
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Globa
191263	28/4/2007	13:07:00	?		?	?
193623	30/4/2007	4:27:00	?		?	?
193619	30/4/2007	4:23:00	?		?	?
193620	30/4/2007	4:24:00	?		?	?
191975	29/4/2007	0:59:00	?		?	?
...
190742	28/4/2007	4:26:00	?		?	?
190499	28/4/2007	0:23:00	?		?	?
190498	28/4/2007	0:22:00	?		?	?
191202	28/4/2007	12:06:00	?		?	?
190497	28/4/2007	0:21:00	?		?	?

4000 rows × 9 columns

```
#sort acc. to index in dec order as asc =false
df.sort_index(ascending=False)
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Glo
1048574	13/12/2008	21:38:00	0.422		0.078	242.61
1048573	13/12/2008	21:37:00	0.422		0.078	242.56
1048572	13/12/2008	21:36:00	0.422		0.076	241.73
1048571	13/12/2008	21:35:00	0.424		0.076	242.1
1048570	13/12/2008	21:34:00	0.426		0.076	242.27
...
4	16/12/2006	17:28:00	3.666		0.528	235.68
3	16/12/2006	17:27:00	5.388		0.502	233.74
2	16/12/2006	17:26:00	5.374		0.498	233.29
1	16/12/2006	17:25:00	5.36		0.436	233.63
0	16/12/2006	17:24:00	4.216		0.418	234.84

1048575 rows × 9 columns

```
df1=df.drop(columns=['Global_active_power'],axis=1)
df1.head()
```

	Date	Time	Global_reactive_power	Voltage	Global_intensity	Sub_metering
0	16/12/2006	17:24:00		0.418	234.84	18.4
1	16/12/2006	17:25:00		0.436	233.63	23
2	16/12/2006	17:26:00		0.498	233.29	23
3	16/12/2006	17:27:00		0.502	233.74	23
4	16/12/2006	17:28:00		0.528	235.68	15.8

```
df2=df.drop([1,3],axis=0) #axis=0 for row
df2.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_ir
0	16/12/2006	17:24:00		4.216	0.418	234.84
2	16/12/2006	17:26:00		5.374	0.498	233.29
4	16/12/2006	17:28:00		3.666	0.528	235.68
5	16/12/2006	17:29:00		3.52	0.522	235.02
6	16/12/2006	17:30:00		3.702	0.52	235.09

```
df.sample(10) #selesct random 10 rows and print it
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Glot
552835	4/1/2008	15:19:00		4.92	0	239.92
619653	20/2/2008	0:57:00		0.24	0	241.15
178361	19/4/2007	14:05:00		2.572	0.2	236.81
851393	29/7/2008	23:17:00		0.45	0.132	241.61
856680	2/8/2008	15:24:00		0.144	0	243.04
911356	9/9/2008	14:40:00		0.188	0	241.06
161892	8/4/2007	3:36:00		2.45	0.216	240
971418	21/10/2008	7:42:00		2.05	0.246	237.05
265580	19/6/2007	3:44:00		0.16	0.1	239.54
903605	4/9/2008	5:29:00		0.31	0.19	238.84

```
print(df.to_string())
```

→ IOPub data rate exceeded.
The notebook server will temporarily stop sending output
to the client in order to avoid crashing it.
To change this limit, set the config variable
`--NotebookApp.iopub_data_rate_limit`.

```
Current values:
NotebookApp.iopub_data_rate_limit=1000000.0 (bytes/sec)
NotebookApp.rate_limit_window=3.0 (secs)
```

```
# myvar=pd.DataFrame(df)
# myvar.head()
```

```
print(df[0])
```

```
234.84
```

```
df
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Glc
0	16/12/2006	17:24:00	4.216		0.418	234.84
1	16/12/2006	17:25:00	5.36		0.436	233.63
2	16/12/2006	17:26:00	5.374		0.498	233.29
3	16/12/2006	17:27:00	5.388		0.502	233.74
4	16/12/2006	17:28:00	3.666		0.528	235.68
...
1048570	13/12/2008	21:34:00	0.426		0.076	242.27
1048571	13/12/2008	21:35:00	0.424		0.076	242.1
1048572	13/12/2008	21:36:00	0.422		0.076	241.73
1048573	13/12/2008	21:37:00	0.422		0.078	242.56
1048574	13/12/2008	21:38:00	0.422		0.078	242.61

1048575 rows × 9 columns

```
#to remove empty cells
df.dropna(inplace=True) #to makes changes perm.
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_ir
0	16/12/2006	17:24:00	4.216		0.418	234.84
1	16/12/2006	17:25:00	5.36		0.436	233.63
2	16/12/2006	17:26:00	5.374		0.498	233.29
3	16/12/2006	17:27:00	5.388		0.502	233.74
4	16/12/2006	17:28:00	3.666		0.528	235.68

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 1044506 entries, 0 to 1048574
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Date            1044506 non-null   object 
 1   Time             1044506 non-null   object 
 2   Global_active_power  1044506 non-null   object 
 3   Global_reactive_power 1044506 non-null   object 
 4   Voltage          1044506 non-null   object 
 5   Global_intensity 1044506 non-null   object 
 6   Sub_metering_1    1044506 non-null   object 
 7   Sub_metering_2    1044506 non-null   object 
 8   Sub_metering_3    1044506 non-null   float64
dtypes: float64(1), object(8)
memory usage: 79.7+ MB
```

```
#replace empty values
df.fillna(145,inplace=True)
df.head(100)
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_i
0	16/12/2006	17:24:00	4.216	0.418	234.84	
1	16/12/2006	17:25:00	5.36	0.436	233.63	
2	16/12/2006	17:26:00	5.374	0.498	233.29	
3	16/12/2006	17:27:00	5.388	0.502	233.74	
4	16/12/2006	17:28:00	3.666	0.528	235.68	
...
95	16/12/2006	18:59:00	4.224	0.09	231.96	
96	16/12/2006	19:00:00	4.07	0.088	231.99	
97	16/12/2006	19:01:00	3.612	0.09	232.36	
98	16/12/2006	19:02:00	3.458	0.09	232.71	
99	16/12/2006	19:03:00	3.434	0.09	232.01	

100 rows × 9 columns

```
#replace null values for a specified column
df['Sub_metering_3'].fillna(145,inplace=True)
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_ir
0	16/12/2006	17:24:00	4.216	0.418	234.84	
1	16/12/2006	17:25:00	5.36	0.436	233.63	
2	16/12/2006	17:26:00	5.374	0.498	233.29	
3	16/12/2006	17:27:00	5.388	0.502	233.74	
4	16/12/2006	17:28:00	3.666	0.528	235.68	

```
#pandas use mean median and mode to replace the empty values
x=df['Sub_metering_3'].mean()
df['Sub_metering_3'].fillna(x,inplace=True)
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_ir
0	16/12/2006	17:24:00	4.216	0.418	234.84	
1	16/12/2006	17:25:00	5.36	0.436	233.63	
2	16/12/2006	17:26:00	5.374	0.498	233.29	
3	16/12/2006	17:27:00	5.388	0.502	233.74	
4	16/12/2006	17:28:00	3.666	0.528	235.68	

```
x=df['Sub_metering_3'].median()
df['Sub_metering_3'].fillna(x,inplace=True)
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_ir
0	16/12/2006	17:24:00	4.216	0.418	234.84	
1	16/12/2006	17:25:00	5.36	0.436	233.63	
2	16/12/2006	17:26:00	5.374	0.498	233.29	
3	16/12/2006	17:27:00	5.388	0.502	233.74	
4	16/12/2006	17:28:00	3.666	0.528	235.68	

```
x=df['Sub_metering_3'].mode()
df['Sub_metering_3'].fillna(x,inplace=True)
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_ir
0	16/12/2006	17:24:00	4.216	0.418	234.84	
1	16/12/2006	17:25:00	5.36	0.436	233.63	
2	16/12/2006	17:26:00	5.374	0.498	233.29	
3	16/12/2006	17:27:00	5.388	0.502	233.74	
4	16/12/2006	17:28:00	3.666	0.528	235.68	

```
#converting date column into dates
df.dtypes
```

→	Date	object
	Time	object
	Global_active_power	object
	Global_reactive_power	object
	Voltage	object
	Global_intensity	object
	Sub_metering_1	object
	Sub_metering_2	object
	Sub_metering_3	float64
	dtype:	object

```
df['Date']=pd.to_datetime(df['Date'])
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intens
0	2006-12-16	17:24:00	4.216	0.418	234.84	1
1	2006-12-16	17:25:00	5.36	0.436	233.63	
2	2006-12-16	17:26:00	5.374	0.498	233.29	

```
df.dtypes
```

→	Date	datetime64[ns]
	Time	object
	Global_active_power	object
	Global_reactive_power	object
	Voltage	object
	Global_intensity	object
	Sub_metering_1	object
	Sub_metering_2	object
	Sub_metering_3	float64
	dtype:	object

```
df['Voltage']=pd.to_numeric(df['Voltage'])
df.head()
df.dtypes
```

→	Date	datetime64[ns]
	Time	object
	Global_active_power	object
	Global_reactive_power	object
	Voltage	float64
	Global_intensity	object
	Sub_metering_1	object
	Sub_metering_2	object
	Sub_metering_3	float64
	dtype:	object

Start coding or [generate](#) with AI.

```
#replacing wrong values with right
df.loc[1.0,'Global_active_power']=45
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity
0.0	2006-12-16	17:24:00		4.216		0.418
1.0	2006-12-16	17:25:00		45		0.436
2.0	2006-12-16	17:26:00		5.374		0.498

```
#finding and removing duplicates true if duplicates false if no duplicates
print(df.duplicated())
```

0.00	False
1.00	False
2.00	False
3.00	False
4.00	False
...	
1048571.00	False
1048572.00	False
1048573.00	False
1048574.00	False
5.36	False

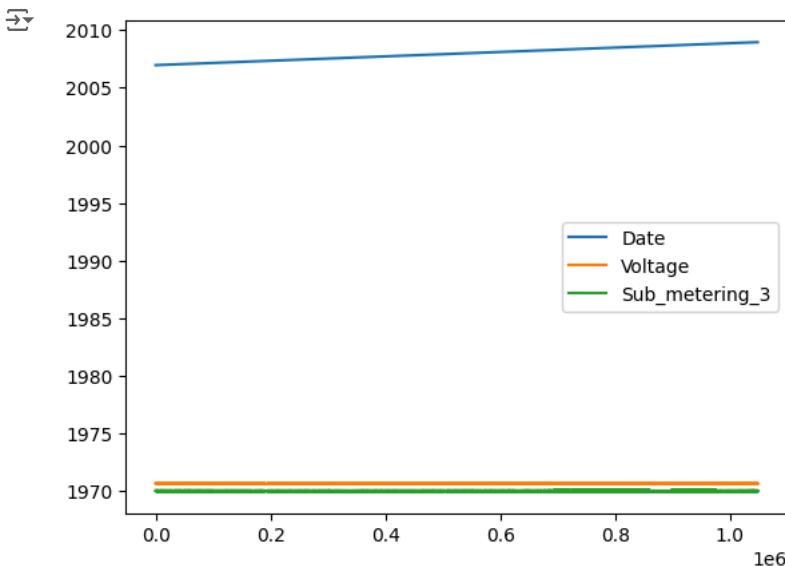
Length: 1044507, dtype: bool

```
#removing duplicates
df.drop_duplicates(inplace=True)
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power	Voltage	Global_intensity
0.0	2006-12-16	17:24:00		4.216		0.418
1.0	2006-12-16	17:25:00		45		0.436
2.0	2006-12-16	17:26:00		5.374		0.498

```
import pandas as pd
import matplotlib.pyplot as plt
df.plot()

plt.show()
```



```
df.head()
df.dtypes
```

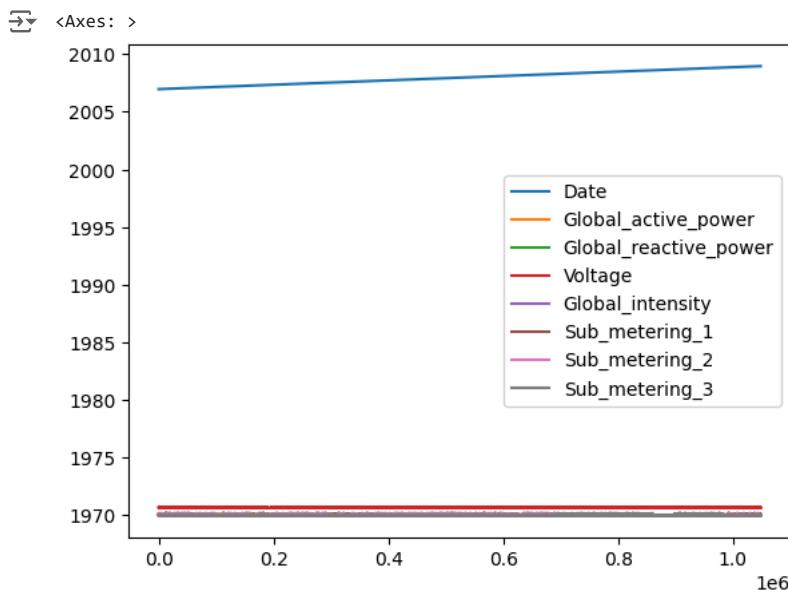
Date	datetime64[ns]
Time	object
Global_active_power	object
Global_reactive_power	object

```
Voltage           float64
Global_intensity    object
Sub_metering_1        object
Sub_metering_2        object
Sub_metering_3        float64
dtype: object
```

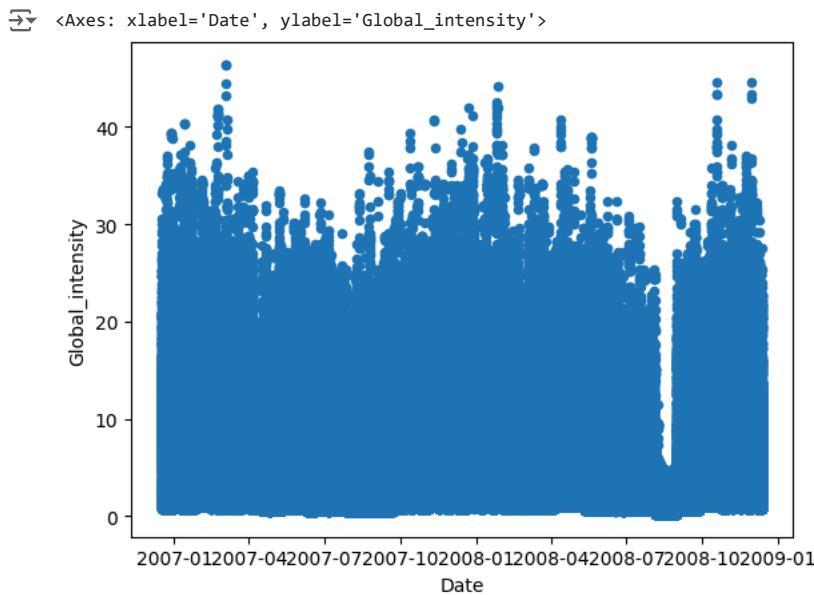
```
df['Sub_metering_2']=pd.to_numeric(df['Sub_metering_2'])
df['Sub_metering_1']=pd.to_numeric(df['Sub_metering_1'])
df['Global_active_power']=pd.to_numeric(df['Global_active_power'])
df['Global_reactive_power']=pd.to_numeric(df['Global_reactive_power'])
df['Global_intensity']=pd.to_numeric(df['Global_intensity'])
df.dtypes
```

```
→ Date          datetime64[ns]
Time           object
Global_active_power   float64
Global_reactive_power  float64
Voltage         float64
Global_intensity    float64
Sub_metering_1      float64
Sub_metering_2      float64
Sub_metering_3      float64
dtype: object
```

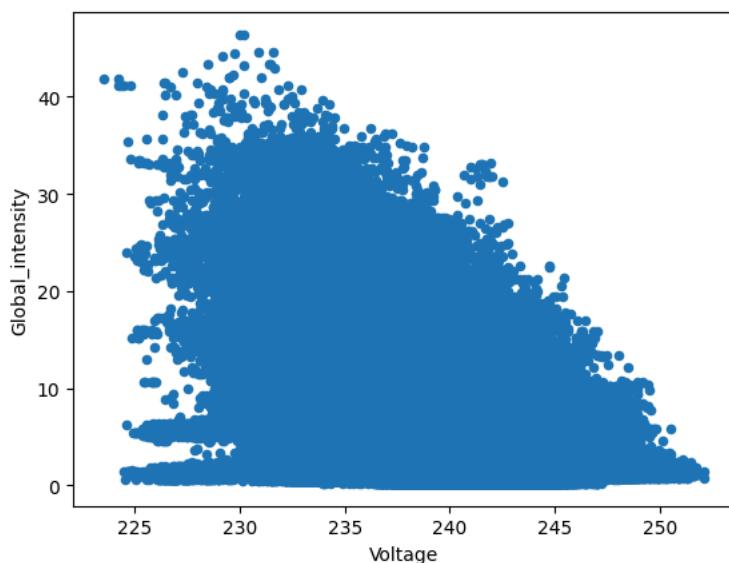
```
import pandas as pd
import matplotlib.pyplot as plt
df.plot()
```



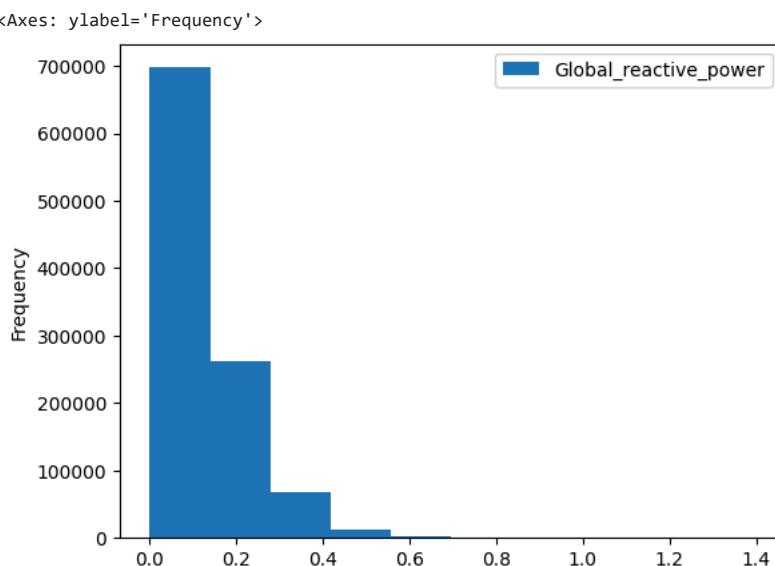
```
df.plot(kind='scatter',x='Date',y='Global_intensity')
```



```
df.plot(kind='scatter',x='Voltage',y='Global_intensity')
plt.show()
```



```
df.plot(kind='hist',y='Global_reactive_power',x='Voltage')
```



```
df.shape
```

(1044507, 9)

```
#no.of rows in series if 2D X with no. of columns
df.size
```

9400563

```
df.ndim
```

2

```
df.describe() #retruns mean median mode Q1 and Q3,etc
```

	Date	Global_active_power	Global_reactive_power	Voltage	Global_intensity
count	1044506	1.044507e+06	1.044506e+06	1.044506e+06	1.044506e+06
mean	2007-12-16 03:01:45.445445120	1.108282e+00	1.182732e-01	2.399598e+02	1.108282e+00
std	1.108282e+00	1.182732e-01	2.399598e+02	1.108282e+00	1.108282e+00
min	2007-12-16 03:01:45.445445120	1.108282e+00	1.182732e-01	2.399598e+02	1.108282e+00
25%	2007-12-16 03:01:45.445445120	1.108282e+00	1.182732e-01	2.399598e+02	1.108282e+00
50%	2007-12-16 03:01:45.445445120	1.108282e+00	1.182732e-01	2.399598e+02	1.108282e+00
75%	2007-12-16 03:01:45.445445120	1.108282e+00	1.182732e-01	2.399598e+02	1.108282e+00
max	2007-12-16 03:01:45.445445120	1.108282e+00	1.182732e-01	2.399598e+02	1.108282e+00